



LUPINE SCIENCE • RESEARCH BOOKLET • V1

The Savings Stack

Seven ways atomistic simulation learned to do less — the measured economics of sharing the expensive part, and the commons that makes every team faster.



July 21, 2026

Free download • lupine.science

Full detail • library.lupine.science

Evidence before claim, in a portable form.

This booklet is the public edition of the Savings Stack v1 release: ten citation-verified chapters on compute-savings techniques in atomistic simulation, one machine-readable measurement of cross-model anchor sharing, and the identifier audit that checked every citation. It is a research release note, not a peer-reviewed publication — every number traces to a committed record you can download, hash, and recompute.

The one-sentence version. The field independently invented seven layers of compute savings, every layer ships without a correctness certificate, and our reviewed measurement says verified gating shared across models delivers 72.4% fewer DFT evaluations.

CONTENTS

The argument	03
The cost wall	04
The stack, drawn to scale	05
Seven layers, one hole	06
The big number	07
The chart	08
The measurement, stated plainly	09
The theorem commons	10
How the commons compounds	11
The demand side already exists	12
What we do not claim	13
The citation audit	14
Data & reproduction	15

Savings Stack v1 · Released 2026-07-21 · Data record z1-union-anchor-economics.json · sha256 22c0ba51d0c798c2... (full hashes in the download pack) · Structured data ODbL-1.0 · Booklet © Lupine Science, content license as per repository · Not peer-reviewed; identifiers and arithmetic verified as recorded on page 14.

01 • THE ARGUMENT

Atomistic simulation spent forty years making FLOPs cheaper — and the last ten discovering that the real enemy is not the price of a FLOP but the *number of times you must call the expensive oracle at all*.

Every successful savings technology of the machine-learning era — active learning, surrogate transition-state search, Δ -learning, abstention gates — is an evaluation-count reduction with an unpriced correctness risk attached.

Our measured contribution: correctness-gated evaluation, *shared across models*, is itself a savings technology — and it scales sub-linearly. If that property generalizes, a commons in which teams contribute theorems and anchors back is a machine where we all get faster together, at no marginal cost to anyone.

That is the whole argument. The rest of this booklet is the evidence.

02 • THE COST WALL

Compute got cheap. Answers did not.

Atomistic campaigns remain constrained by expensive reference evaluations, but the prior cost, electricity, throughput, and scaling figures on this page were derived or inherited from literature digests that do not yet have a reviewed public source contract.

FAIL-CLOSED BOUNDARY

Quantitative compute economics are withheld. The qualitative question remains: who reduces reference work, and under what correctness boundary?

The underlying backdrop chapters are preserved outside the deployed public tree pending renewed claim review.

03 • THE STACK, DRAWN TO SCALE

Seven layers of savings.
One crack through all of them.

*Seven technical families.
One shared need for a
correctness certificate.*

QUANTITATIVE LITERATURE SAVINGS ARE WITHHELD PENDING RENEWED CLAIM
REVIEW.

Each layer is measured, replicated, and citable — the strongest verified number is drawn on its stone. And each layer's own literature names the same hole: the trigger, the variance, the committee estimates *where the model is probably wrong*. None can say *where physics is definitely violated*.

The cited literature digests remain preserved as review evidence outside the public deployment.

04 · SEVEN LAYERS, ONE HOLE

What each layer saves — and what each one cannot say.

01	Surrogate transition-state search Uses a learned local surface to reduce reference work.	THE HOLE GP variance is a sampling-density signal, not an accuracy bound.
02	Active learning Chooses reference labels from model disagreement.	THE HOLE A committee can agree and be wrong under extrapolation.
03	Multi-fidelity transfer Combines lower- and higher-fidelity evidence.	THE HOLE Transfer can soften the target surface out of distribution.
04	Abstention Routes uncertain cases back to a reference method.	THE HOLE Cross-model sharing requires a reviewed correctness boundary.
05	Systems-level DFT acceleration Reduces work inside each reference evaluation.	THE HOLE A faster evaluation can still fail silently.
06	Electronic-structure surrogates Predicts electronic quantities without a full reference solve.	THE HOLE Predicted Hamiltonians do not guarantee variationally reliable forces.
07	Path and sampling algorithms Changes how candidate pathways are explored.	THE HOLE Cheap exploration can skip the transition-state region.

The detailed rows and their explicit [UNVERIFIED] flags are preserved outside public deployment pending renewed review.

05 • THE BIG NUMBER

72.4%


fewer DFT evaluations when four models *share their anchors* instead of each paying for its own.

The reviewed result is bounded to 29 analyzable diffusion paths on the locked Z1 panel. It was measured from recorded campaign artifacts; no new DFT was run to produce it. Detailed arithmetic remains in the hashable evidence record rather than being restated as substitute public economics.

z1-union-anchor-economics.json · sha256 22c0ba51d0c798c2... · recomputed 2026-07-21 from the Round-4 float64 campaign artifacts under the frozen anchor rule of the preregistered sparse-DFT pilot.

06 • THE CHART

Separate guidance. Shared reference anchors.



72.4% fewer DFT evaluations on the reviewed panel.

Models largely nominate the same images, so one reference evaluation can serve multiple guidance routes. The public booklet states only the frozen reviewed result; detailed supporting arithmetic remains withheld pending ledger verification.

The detailed analysis record is preserved outside the public pack pending ledger verification.

07 · THE MEASUREMENT, STATED PLAINLY

One panel, four models, and an honest denominator.

The locked Z1 panel holds 30 diffusion paths. Path 14 (mp-756912_1_1_1_0_0) failed CI-NEB convergence under the frozen protocol in *all four* model artifacts, so no profile exists to guide anchors — the analyzable denominator is 29 of 30, on record, with per-model failure records in the JSON. Three more paths carry partial coverage; agreement there is reported under two conventions. Nothing was imputed.

APPROVED PUBLIC RESULT

72.4% fewer DFT evaluations. Bounded to the 29 analyzable paths on the locked panel; recorded artifacts only, with no new DFT run for the review.

DO THE MODELS EVEN AGREE? PREDICTED SADDLE/EXTREMA COINCIDENCE OVER 29 PATHS

STATISTIC	AVAILABLE - MODELS	ALL - FOUR - REQUIRED
Saddle (argmax) exact	23/29	20/29
Saddle within ±1 image	26/29	23/29
Both extrema exact	10/29	9/29
Both extrema within ±1	12/29	11/29

The anchor rule is frozen: per model and path, endpoints plus the model-predicted minimum and a window around the model-predicted maximum. The naive total pays every model’s set; the union total pays one evaluation per unique image per path. The protocol itself was validated before any DFT ran — sparse-protocol barrier MAE of 1.2–9.4 meV at ~7 anchors per path against ~50–100 dense evaluations — and one anchor point has since landed on real metal: 32.2 meV against a frozen 40 meV gate.

PUBLIC ECONOMICS BOUNDARY

Ratios, rounded substitutes, curve comparisons, and compounded estimates remain outside approved public copy. The downloadable record preserves the underlying measurements for review without turning them into additional public claims.

The public claim contract is approved-public-economics.json in the pack. Detailed arithmetic and per-path tables are withheld from deployment pending a newly reviewed source contract.

08 • THE THEOREM COMMONS

*“Every team that joins
makes every other team
faster.”*

THE THEOREM COMMONS, IN ONE SENTENCE

09 • HOW THE COMMONS COMPOUNDS

Theorems are non-rival goods with zero marginal sharing cost.

A machine-checked physical-law theorem — the curvature signature at a first-order saddle, energy–force consistency, symmetry constraints, boundary conditions — is contributed once and then gates every simulation for everyone, forever. Unlike training data, *a theorem leaks nothing about the contributor's chemistry or commercial targets*. It is the shareable residue of work a serious team does anyway: every failed simulation is a candidate theorem, and formalizing your own failure converts your most expensive knowledge into a permanent asset for a commons that pays you back in everyone else's theorems.

- 1 Run your campaign on the commons. Your model's known failure modes are caught by existing theorems — you save evaluations immediately.
- 2 Your novel failures get formalized. Each new failure mode becomes a new contributed theorem.
- 3 Every member's gates get sharper. False-accept and false-reject rates both fall, network-wide.
- 4 Savings compound. Evaluations saved grow with membership — gate quality is a monotone function of it.

The commons has two tiers, and members choose their exposure. Theorems are zero-leak by construction. Anchor libraries — a content-addressed store mapping structure-hash to DFT energy and forces — are opt-in per project: one team's anchor validates every other team's model on the same path, under the same arithmetic measured on page 07.

The reviewed union-anchor result — 72.4% fewer DFT evaluations — is the measured instance of the evaluation-side commons. Theorem layer: the machine-generated Lean inventory in the open library, zero sorry.

10 • THE DEMAND SIDE ALREADY EXISTS

Seven digests, working independently, came back with the same hole. Read the literature's own concessions:

Learned triggers are *sufficient-not-necessary*— the committee can agree and be wrong (DP-GEN). GP variance *underestimates* error under strong extrapolation (FLARE). Ensemble σ *decorrelates* from true error. Fast exploration *skips* the transition-state region (OneOPES). String methods land on *second-order* saddles at measured rates. ML functionals fail *silently* (DM21).

Every one of those failure modes is a *physical-law violation detectable without learning anything*. A commons of formal theorems is precisely the supply for a demand the field has been circling for five years.

11 • RESTRAINT, AS A POLICY

What we do not claim.

01 Not generality beyond the measured basis.

Our measurement is one 30-path barrier panel, four universal MLIPs, one DFT engine (GPAW, frozen PBE/fd settings), one chemistry family. The union scaling curve is four points deep.

02 Not unreviewed stacking arithmetic as public economics.

Ratios, rounded substitutes, and compounded estimates are withheld from public copy unless a new source contract is reviewed.

03 Not victory over the point solutions.

For a single isolated saddle search with a cheap oracle, local GP surrogates win outright. For one stable, high-volume chemistry, fine-tuning is the right tool. Our wedge is amortized, multi-path, multi-model workloads — panels, networks, screening campaigns.

04 Not blanket validation of every literature number.

The digests retain their [UNVERIFIED] claim flags and per-source access levels. The audit verified every identifier — it does not upgrade abstract-only evidence or independently reproduce reported savings factors.

05 Not peer review.

This is a research release with committed, hashable records — an invitation to inspect and recompute, not a claim of external validation.

Restraint is the brand: a trust layer that overclaims is not one.

12 • THE CITATION AUDIT

Every identifier, checked against its registry.

All 163 unique identifiers extracted from the nine digests were verified against the official registries — arXiv IDs against the arXiv Atom API, DOI strings against Crossref Works metadata — with returned titles reviewed against the citing rows. That review caught one live but unrelated DOI hidden in a link target.

115/115

arXiv identifiers verified

48/48DOI strings verified through
Crossref

0

unresolved identifiers

Corrections applied on the record. The Chen et al. multi-fidelity citation was moved from a nonexistent DOI to Nature Computational Science 1, 46–53 (10.1038/s43588-020-00002-x). The Nandi 2021 link target — which pointed at a live but unrelated paper — was rebound to the citation text’s own DOI (10.1063/5.0038301). And the VASP MLFF venue was reconciled: PRB 100, 014105 / arXiv:1904.12961 is the melting-point paper behind the >99%-bypass claim; PRL 122, 225701 is the distinct earlier hybrid-perovskite demonstration.

WHAT THE AUDIT DOES NOT DO

Identifier verification does not upgrade abstract-only evidence, does not remove each digest’s [UNVERIFIED] quantitative-claim flags, and does not independently reproduce reported savings factors. It says the citations resolve to what we say they are — no more, and no less.

Full audit record: citation-verification-2026-07-21.md in the pack, with the per-identifier machine-readable artifact and the reproducible checker named there.

Download it. Hash it. Recompute it.

The Savings Stack v1 public pack is deliberately narrow: it ships only the reviewed claim contract, this withholding notice, and a sha256 manifest. The larger evidence pack is preserved outside the deployed public tree until its quantitative claims receive a new source contract.

approved-public-economics.json	the machine-readable reviewed claim contract
README.md	the fail-closed publication boundary
MANIFEST.sha256	hashes over both public files

```
$ sha256sum -c MANIFEST.sha256      # every public file: OK
```

Pack: [lupine.science/data/savings-stack-v1/](#) · Booklet: [lupine.science/booklets/the-savings-stack.pdf](#) · Full scientific detail: [library.lupine.science](#), group “savings-stack”.



Evidence before claim.

The Savings Stack · v1

Booklet & data pack – lupine.science

Full detail – library.lupine.science · group “savings-stack”

The instrument – lupi.live